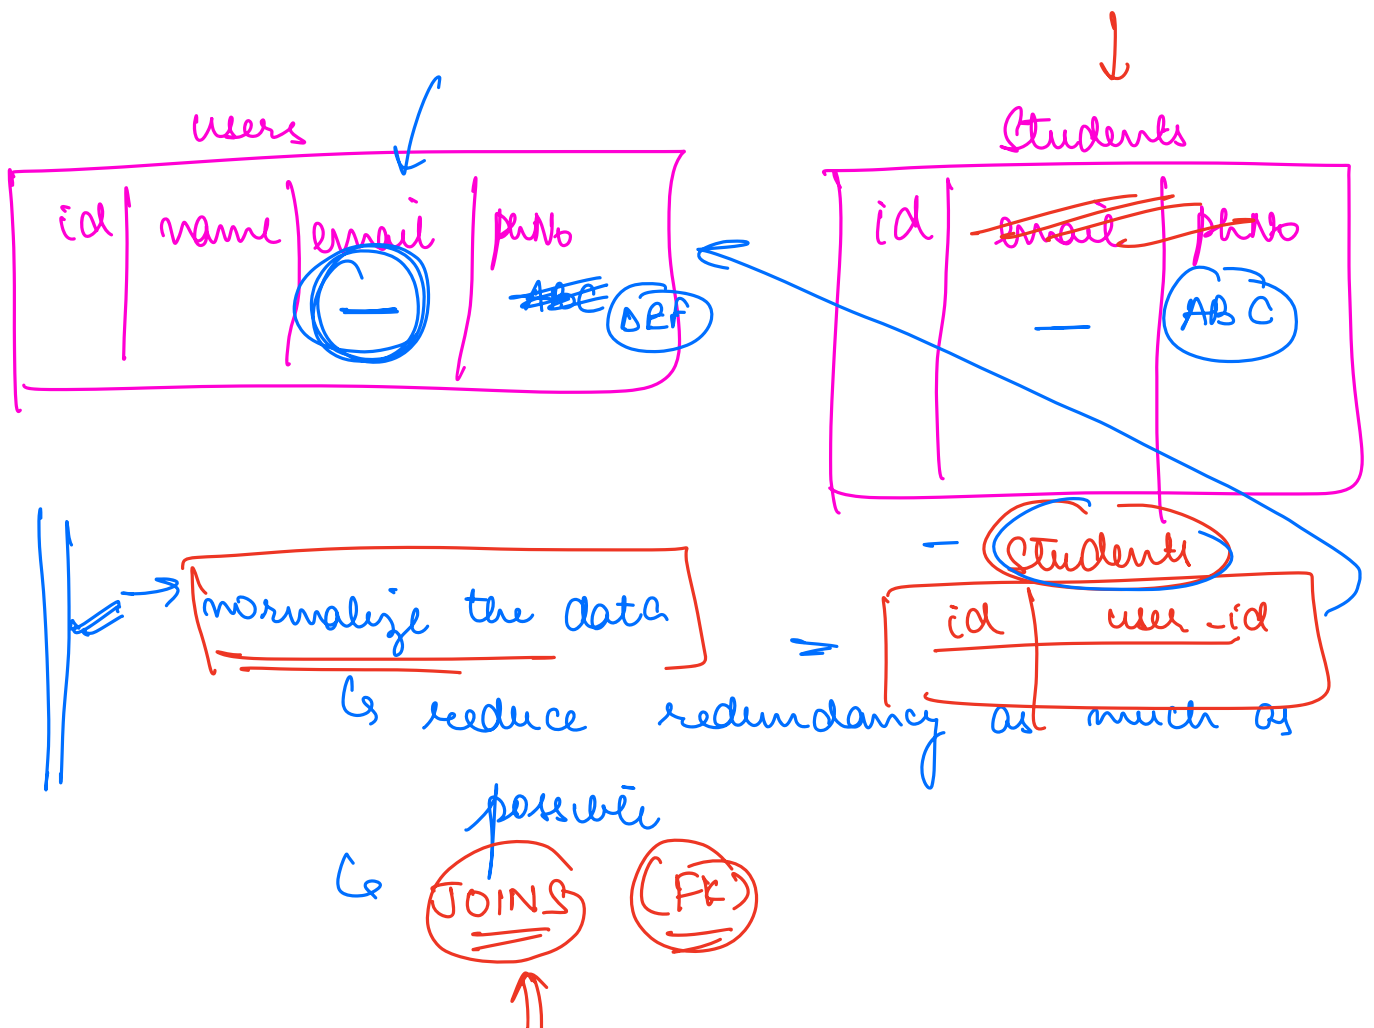


## Agenda

- ① Intro to SQL
  - SQL
  - Shortcomings
  - NOSQL
- ② How to decide Sharding Key




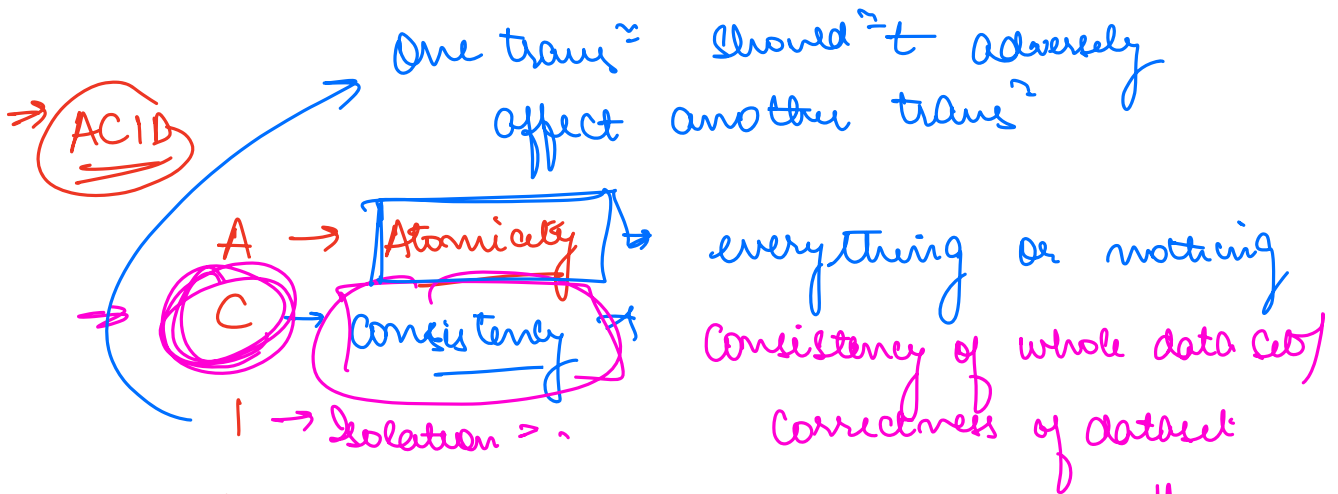
## Features of SQL

→ ACID



⇒ if joins were slow. Then normalizing may

in fact be bad.



Read a → temp  
temp = temp + 1  
Write a ← temp

Atomic Integer a  
a. increment()

SQL → Normalization X

1 → ~~ACID~~ X  
in memory

Most NoSQL

Databases  
actually remove  
one of ACID  
to give their  
properties

Redis: NoSQL Database

→ just write and  
read.

RAM >> SSD >> Hard Disk

Structure

→ Defined Schema

→ defined set of attr

& defined data types

long

varchar(50)

| id | name | email |
|----|------|-------|
|    |      |       |

Ecom Website



products

| id | title    | Category | price  | ram    | os    | Storage | Size |
|----|----------|----------|--------|--------|-------|---------|------|
|    | Chet Sir | phuric   | height | author | ishin | genre   |      |
|    |          |          |        |        |       |         |      |

publisher

→ infinite # of attr of products

→ if SQL ⇒ null values



SQL=1

clothes

|  |  |  |  |
|--|--|--|--|
|  |  |  |  |
|--|--|--|--|

pharmy


|  |  |  |
|--|--|--|
|  |  |  |
|--|--|--|

products

| id | title | meta data                 |
|----|-------|---------------------------|
|    |       | { key: 'm',<br>price: ~ } |

⇒ find all

washing machines



25000 tables

⇒ SQL is optimized for large data in less # of tables  
V/S less data in more # tables

clothes of all

Size.

⇒ VVVV slow.

- ① Normalization ✓
- ② ACID ✓
- ③ Structured ✓

Efficiency  
V/S  
those.



# No SQL Database

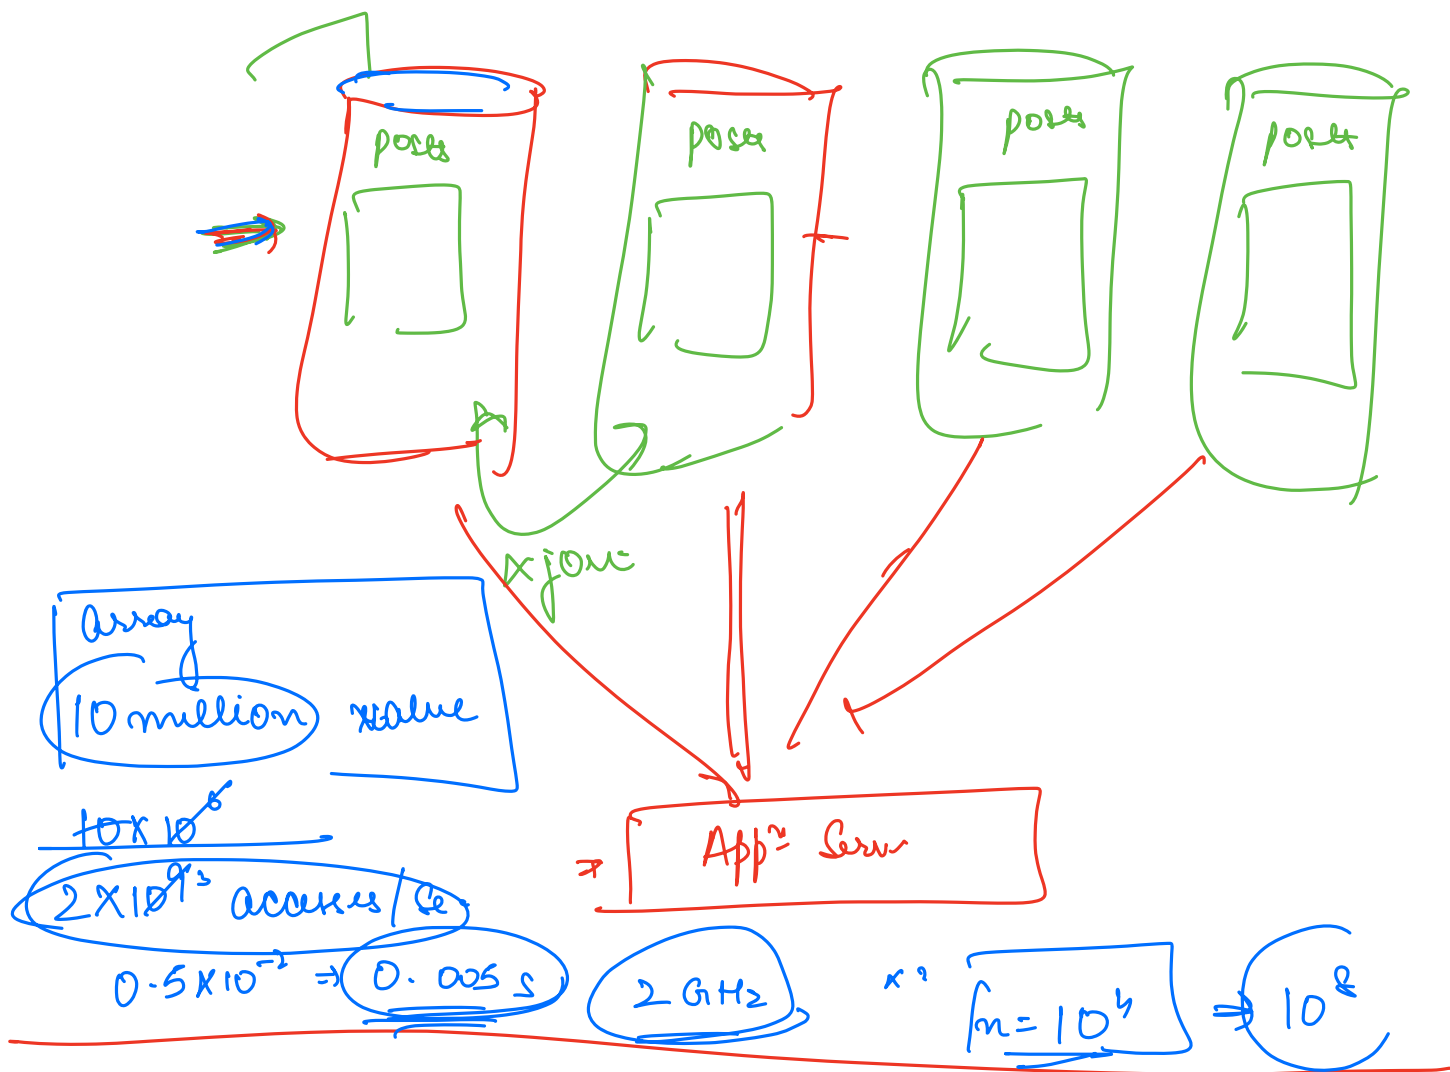
SQL DB work much better than what you want

→ evaluate SQL

→ see if SQL poses some problems basis req

→ if yes: evaluate alternatives





① What happens when size of data is huge.

→ I will choose DB based on type of queries

→ size of data  
→ split data


↳ Ideally split data in a way that SQL DB will work

# → How to choose sharding key


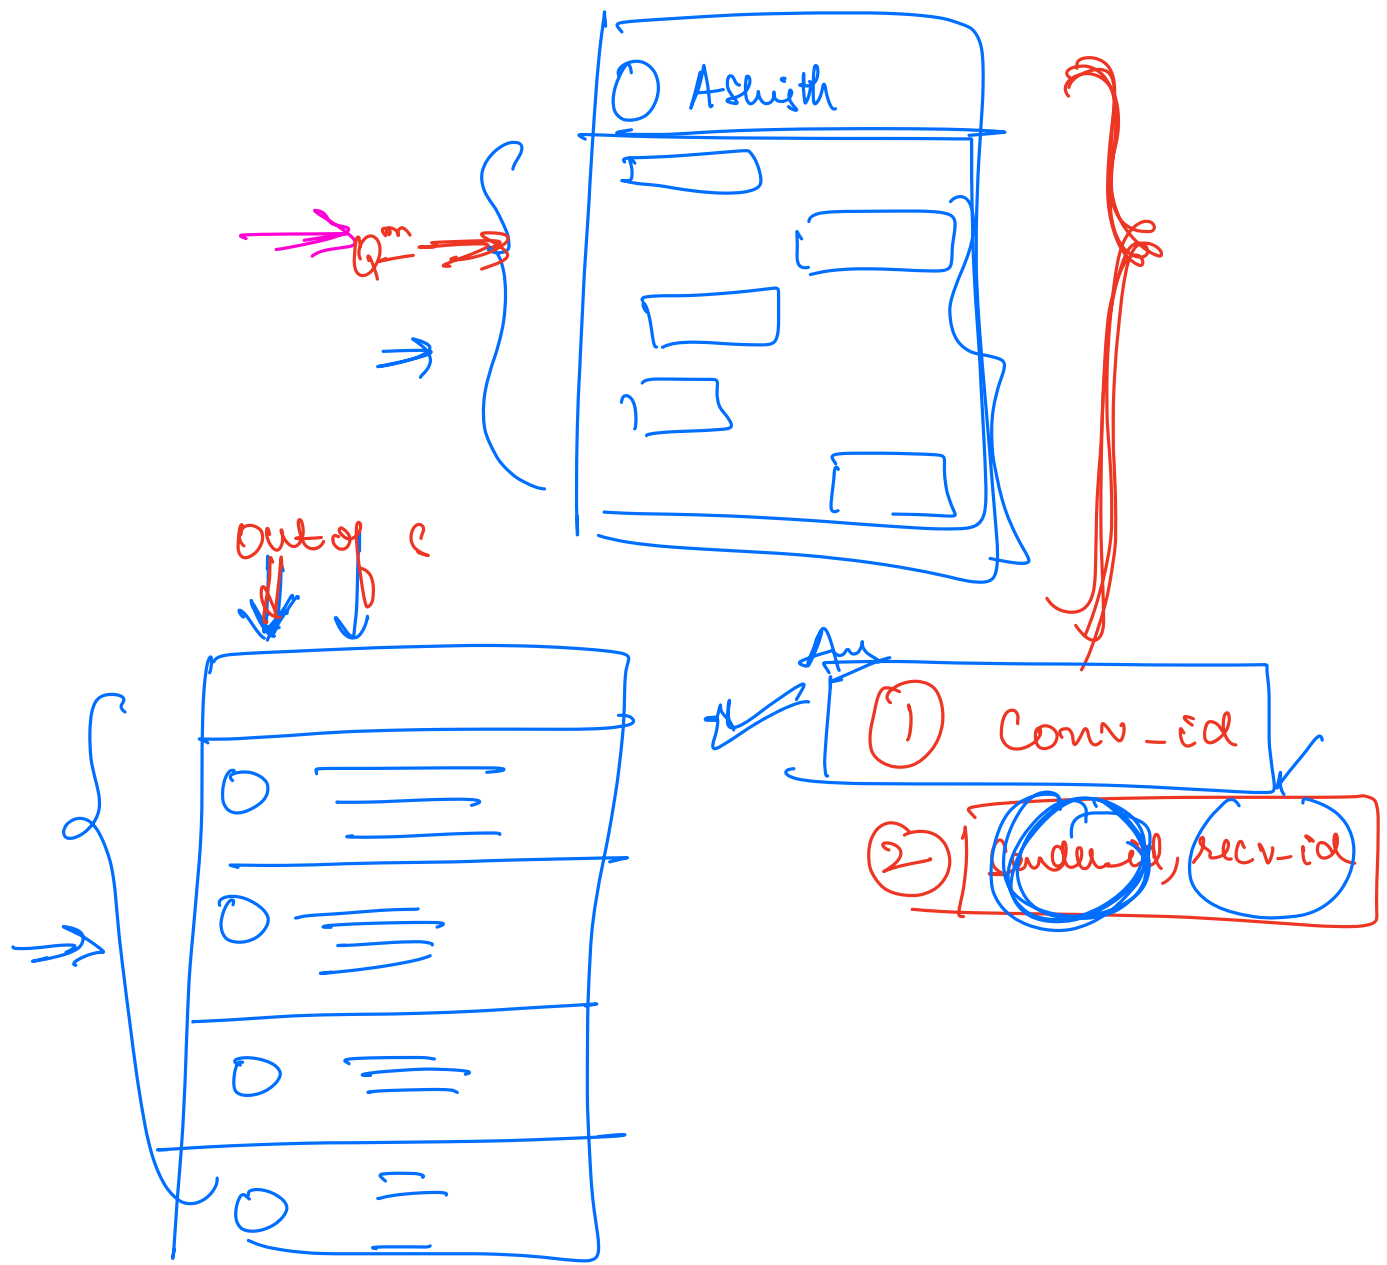
## ① Messenger app<sup>ns</sup>

WhatsApp / FB Messenger

| message |      |           |         |         |        |
|---------|------|-----------|---------|---------|--------|
| id      | time | sender-id | recv-id | Content | status |
|         |      | X         |         |         |        |







Overloaded

→ Select from tickets  
where train-id = 6 & date =

How TO CHOOSE SHARDING KEY

hash(?)

→ creating too small  
sharding key

① makes our query fast  
(ideally only one db  
machine to go to)

② load across shards  
should be uniform

③ don't choose too broad  
sharding key

→

Conversations

| id | user1 | user2 | lastMessageTime |
|----|-------|-------|-----------------|
|    |       |       |                 |
|    |       |       |                 |
|    |       |       |                 |

$h(a, b)$   
→  $h(a \cdot 10^9 + b \cdot 10^4)$

→

messages

| id | conv-id | text | time |
|----|---------|------|------|
|    |         |      |      |
|    |         |      |      |
|    |         |      |      |

# Bank System

## Accounts

| id | Owner Name | branch | address City | type | balance |
|----|------------|--------|--------------|------|---------|
|    |            |        |              |      |         |


## In Account id


Op<sup>m</sup>

- ① get balance of an Account
- ② get transaction history of an account

| transaction |                   |        |             |
|-------------|-------------------|--------|-------------|
| <u>id</u>   | <u>account-id</u> | amount | <u>type</u> |
|             |                   |        |             |

all trans<sup>r</sup> of same acc<sup>t</sup> in same row.





Account id + branch == account.


User

Req → people would want to search cars near to them.

↓

| <del>cars</del> Drivers |      |        |         |          |
|-------------------------|------|--------|---------|----------|
| id                      | name | Car No | city-id | lat long |
|                         |      |        |         |          |

If you land on (lat, long)  $\Rightarrow$  all cars with (lat, long) will be in same machine



$\Rightarrow$  ideally all cars nearby are in same machine

Uber

Google Maps / Uber  $\Rightarrow$  Tinder

Slack

messages

| -id | user-id | group-id | content | time | workspaces |
|-----|---------|----------|---------|------|------------|
|     |         |          |         |      |            |

too abstract  
 $\downarrow$   
 too big

$\Rightarrow$  sharding key should be even that

it doesn't overflow db in itself

ACTC

- ① Same seat shouldn't be booked twice
- ② I should be able to get all bookings for a particular train to be able to assign

| ticket |                     |                 |           |         |                                    |
|--------|---------------------|-----------------|-----------|---------|------------------------------------|
| id     | # people in booking | date of Journey | start loc | end loc | start pt of train and point of loc |
|        |                     |                 |           |         | train-id                           |


potential

- ① train-id + doj
- ② train-id
- ③ start pt of train → too big
- ④ ~~train-id + ticket-id + # of people~~

= [ train-id + ticket-id ]

⇒ all tickets will have same train-id and


ticket id are in one machine




→ all tickets of  
same train  
to be together

→ train id


(1, 24<sup>th</sup> October)



(2, 24<sup>th</sup> October)




(1, 20<sup>th</sup> Jan)



AND

(2, 20<sup>th</sup> Jan)



$\text{hash}(a, b) \Rightarrow \text{hash}((a \times 10^8) \times 10^7 \text{ } \leftarrow)$

TA  
↓

↓

Select ^  
from tickets  
where train-id = 1  
and doj = 2

⇒ Ideal key to  
get all such  
rows on  
same machine